




# Parallel Reinforcement Learning

## Parallel Design

team07: Sam DePaolo, Michael Kielstra, Manqing Liu, Xiaohan Wu

# Refresher: Reinforcement Learning

- At each time  $t$ , the agent receives current **state**  $S_t$  and **reward**  $R_t$
- The agent then choose **action**  $A_t$
- The action is sent to the environment, which moves to a new state  $S_{t+1}$  and reward  $R_{t+1}$
- Goal is to learn a **policy**  $\pi$  which maximizes the cumulative reward
$$\pi(a, s) = Pr(A_t = a \mid S_t = s)$$



## Sequential Baseline Results

- Learning performance of a single RL agent with a 10-armed bandit
- We use an  $\epsilon$ -greedy policy ( $\epsilon = 0.1$ ) to average 500 different experiments where each contain 100 trials
- $n=10$  arms are created randomly from  $N(1.0, 1.0)$
- As the agent gains more **experience**
  - Reward for each arm approaches the **true mean**
  - The agent is more likely to select the **optimal action**





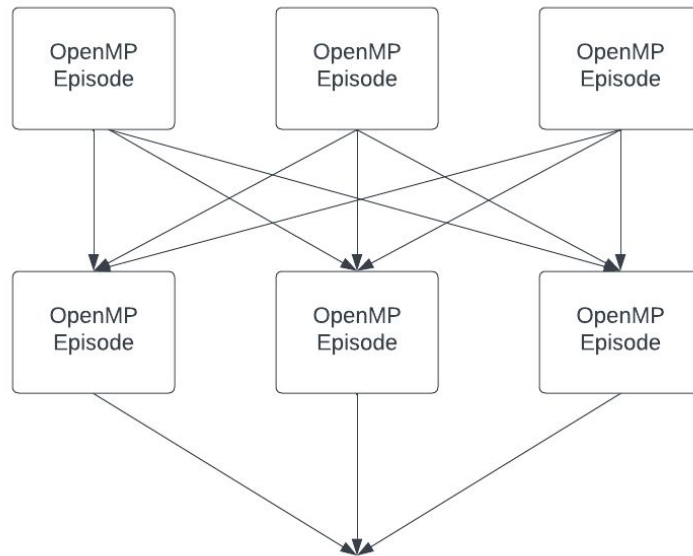
Figure: Single Agent in 10-armed Bandit Task

# Roofline Analysis



# Parallel programming models

- We cannot make the random number generators perform better
- Parallelize both the **agents** and the training within each **episode**
- Parallelize the agents with **MPI**; parallelize the episodes with **OpenMP**





## Plan to implement parallel code

- Flow of computational steps:
  - Define an initial best action
  - Run the 500 episodes with 100 trials each
- Synchronization must occur:
  - Between agents, after a set number of episodes to share the information among agents and optimally update for learning
  - Within an agent, between each episode
- No load imbalance will be expected due to the random and repeated nature of the RL

# Hiding Latency

- Latency will likely occur largely when sharing information between agents:
  - Small size of information makes latencies even more apparent
  - Must determine optimal frequency in which to share between agents to hide latencies
  - Want to send information between agents while calculations continue to happen

